

Understanding Parallelism & Concurrency

(informal introduction)

Mahmoud Parsian

Ph.D. in Computer Science



Why Parallelism?

- **Parallelism and Partitioning are foundations of big data solutions.**
 - MapReduce is based on parallelism and partitioning data
 - Hadoop is based on parallelism and partitioning data
 - Spark is based on parallelism and partitioning data
 - Snowflake is based on parallelism and partitioning data
- **The concept of parallel computing is based on**
 - dividing a large problem into smaller ones, and
 - each of them is carried out by one single processor individually.



Parallelism Example: array addition

```
// a: a[0], a[1], ..., a[n-1]
// b: b[0], b[1], ..., b[n-1]
for (i = 0; i < n; i++) {
    c[i] = a[i] + b[i];
}
```

```
c[0] = a[0] + b[0]
c[1] = a[1] + b[1]
c[2] = a[2] + b[2]
...
c[n-1] = a[n-1] + b[n-1]
```

Sequential:

n steps

one processor

Parallel:

one step

n processors



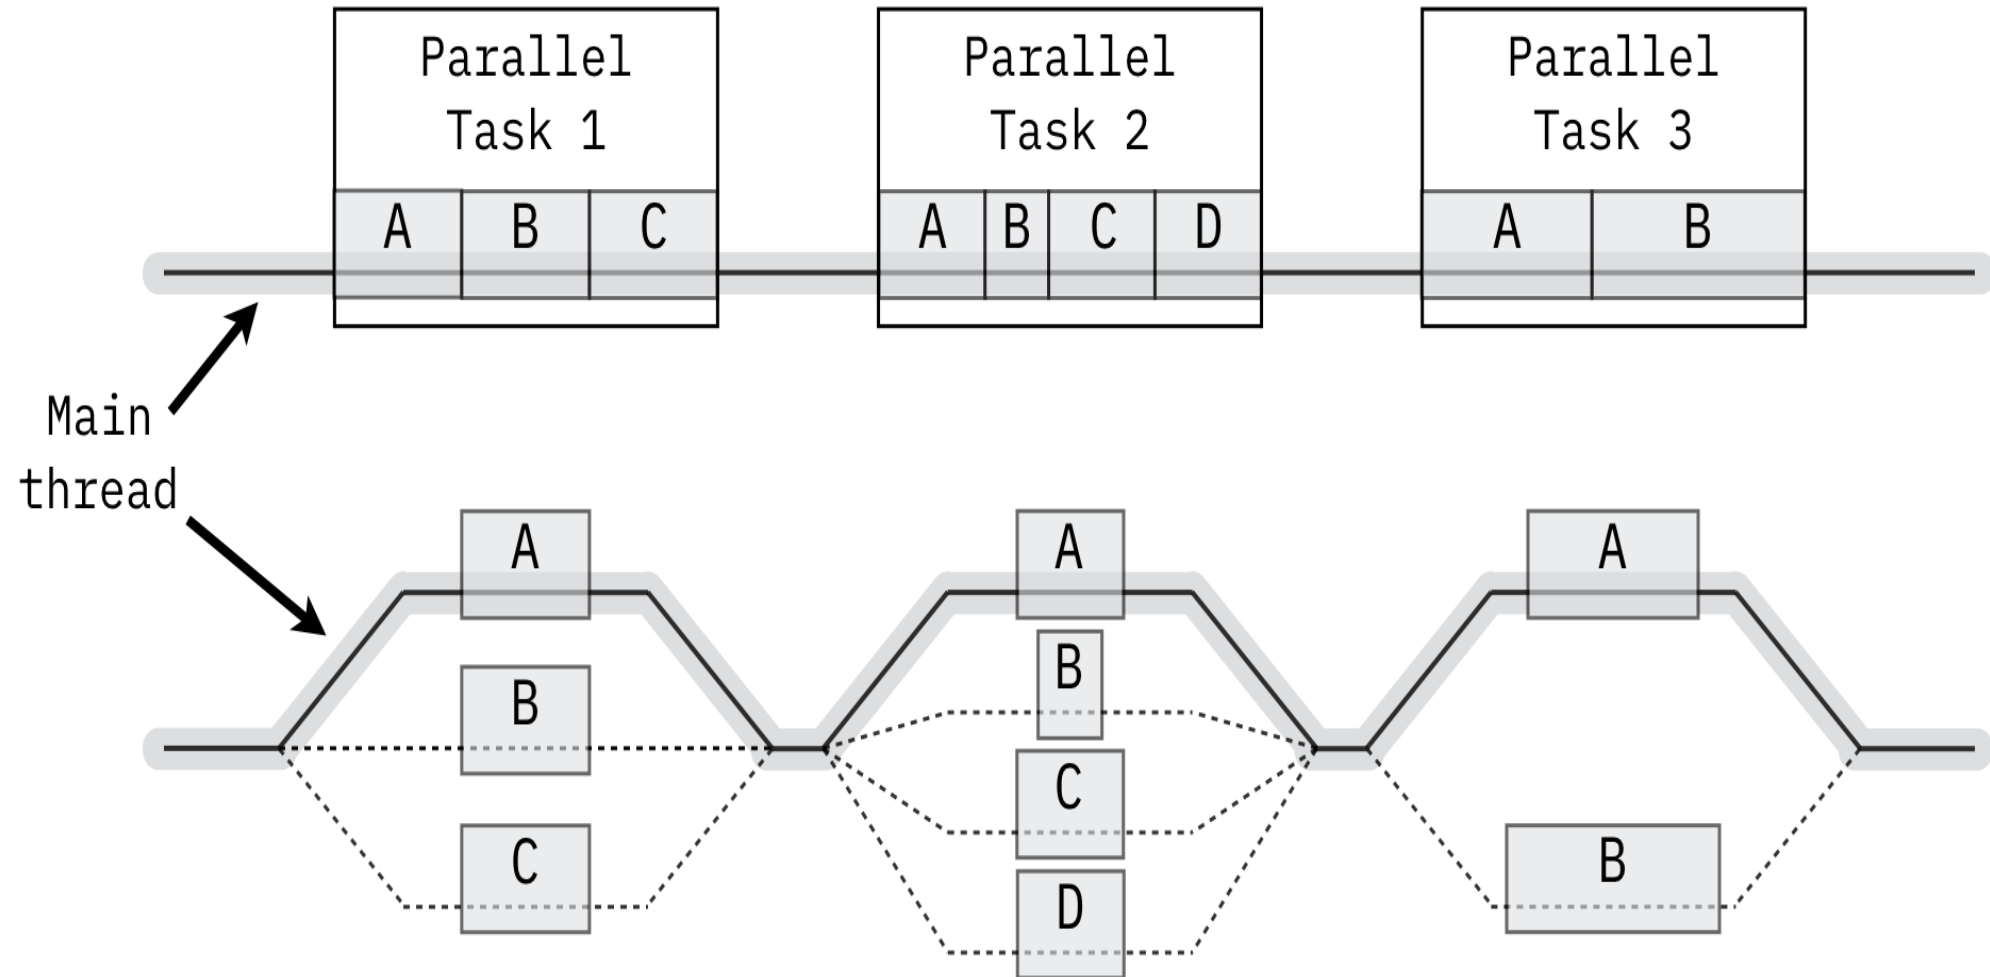
Concurrency & Parallelism (informal definition):

Some definitions:

- The fact of two or more events or circumstances happening or existing at the same time.
- The ability to execute more than one program or task simultaneously.
- Example: "a high level of concurrency is crucial to good performance in a multiuser database system"
 - Multiple reads at a time
 - Multiple writes at a time



Concurrency & Parallelism (informal definition):




Concurrency & Parallelism

Task A = { A1, A2, A3, A4}, Task B = {B1, B2, B3}, Task C = {C1, C2}

Execute Tasks A and B and C (in sequential order):

1. A1
2. A2
3. A3
4. A4
5. B1
6. B2
7. B3
8. C1
9. C2



Concurrency & Parallelism

Task A = { A1, A2, A3, A4} Task B = {B1, B2, B3}, Task C = {C1, C2}

Assume that Tasks (A, B, and C) are independent of each other

Execute Tasks A, B, C: (concurrently)

1. A1 B1 C1 (3 tasks running concurrently)
2. A2 B2 C2 (3 tasks running concurrently)
3. A3 B3 (2 tasks running concurrently)
4. A4 (1 task running)



Concurrency & Parallelism: Maximum Parallelism

Task A = { A1, A2, A3, A4 }

Task B = { B1, B2, B3 },

Task C = { C1, C2 }

If there is NO dependency between any tasks:

Then run all of these 9 tasks concurrently:

A1, A2, A3, A4, B1, B2, B3, C1, C2



Parallelism

Parallelism is basically a type of computation in which many tasks/computations/operations are carried out in parallel.

Sequential Tasks:

Task-1: 20 minutes

Task-2: 20 minutes

Task-3: 25 minutes

Task-4: 25 minutes

=====

Total Elapsed time: 90 minutes

(assuming that there is no dependencies between tasks)



Parallel Tasks:

Task-1, Task-2, Task-3, Task-4

=====

Total Elapsed time: 25 minutes

→ Therefore, Parallelism improves execution time.



Parallelism

What if {Task-3 and Task-4} **depends on output of** {Task-1 and Task-2}

4 Sequential Tasks:

Iteration-1: Task-1 : 20 minutes

Iteration-2: Task-2 : 20 minutes

Iteration-3: Task-3: 25 minutes

Iteration-4: Task-4 : 25 minutes



=====

Total Elapsed time: 90 minutes

(assuming that there is no dependencies between tasks)

Dependencies can be bottlenecks!!!

4 Parallel Tasks:

Iteration-1: Task-1, Task-2 (duration: 20 mins)


Iteration-2: Task-3, Task-4 (duration: 25 mins)



=====

Total Elapsed time: 45 minutes

Parallelism



Birthday Example

- Alex wants to throw a big birthday party.
- He has a list of 1000 items to buy from a Safeway (grocery store).
- Items to buy:
 - item-1: ice-cream
 - item-2: chips
 - item-3: oranges
 - item-4: cup cakes
 - ...
 - item-1000: grapes
- How long will it take for Alex to buy all these 1000 items
 - assuming that he is the only person in grocery store?



What is our Data?

ID	Item to buy
1	ice-cream
2	chips
3	oranges
4	Cup-cakes
...	...
1000	grapes

Birthday Example ...

- **How long will it take for Alex to buy all these 1000 items?**
 - assuming that he is the only person in grocery store
- **It is estimated that on average he needs 14 seconds to find the item and put in the shopping cart.**
- **Therefore, he will spend**
 $1000 \times 14 = 14,000 \text{ seconds} = 234 \text{ minutes}$

BUT, Alex is busy and can not spend 234 minutes for shopping.

What to do?



Processing Data? One Executor: Alex
Buy item-1, then item-2, ..., then item-1000, ...
All is done in sequence: No Parallelism yet.

ID	Item to buy
1	ice-cream
2	chips
3	oranges
4	Cup-cakes
...	...
1000	grapes

Birthday Example ...

But, Alex is **busy** and can not spend **234 minutes** for shopping.

What to do?

- Alex calls 10 of his friends (F1, F2, ..., F10):
 - they will do shopping and deliver to Alex.
- **Therefore, Each friend will buy 100 items**
- **Total time elapsed to buy all items:**
 $100 \times 14 = 1400 \text{ seconds} = 24 \text{ minutes}$
- **THIS IS a BIG IMPROVEMENT: 234 minutes reduced to 24 minutes**
- **Alex is still not satisfied with this plan.**



Processing Data? 10 parallel executors

Each executor operates in parallel & independently

Executor-1: F1

ID	Item to buy
1	ice-cream
2	chips
3	oranges
4	Cup-cakes
...	...
100	grapes

Executor-1: F2

ID	Item to buy
101	...
102	...
103	...
104	...
...	...
200	grapes

...

Executor-10: F-10

ID	Item to buy
901	...
902	...
903	...
904	...
...	...
1000	grapes

Birthday Example ...

But, Alex is busy and can not spend 24 minutes for shopping.

What to do?

- Alex calls 100 of his friends (F1, F2, ..., F100):
they will do shopping and deliver to Alex.
- Therefore, Each friend will buy 10 items
- Total time elapsed to buy all items:
 $10 \times 14 = 140$ seconds = about 3 minutes
- **THIS IS a BIG IMPROVEMENT: 234 minutes reduced to 3 minutes**
- Alex is still not satisfied with this plan.



Processing Data? 100 parallel executors

Each executor operates in parallel & independently

Executor-1: F1

ID	Item to buy
1	ice-cream
2	chips
3	oranges
4	Cup-cakes
...	...
10	...

Executor-2: F2

ID	Item to buy
11	...
12	...
13	...
14	...
...	...
20	...

...

Executor-100: F-100

ID	Item to buy
991	...
992	...
993	...
994	...
...	...
1000	grapes

Birthday Example ...

But, Alex is busy and can not even spend 3 minutes for shopping.

What to do?

- Alex calls **1000** of his friends (F1, F2, ..., F1000):
 - they will do shopping and deliver to Alex.
- **Therefore, Each friend will buy 1 single item**
- **Total time elapsed to buy all items:**
 $1 \times 14 = 14$ seconds
- **THIS IS a HUGE IMPROVEMENT:**
234 minutes reduced to 14 seconds
- **Alex is satisfied with this plan.**



Processing Data? 1000 parallel executors

Each executor operates in parallel & independently

Executor-1: F1

ID	Item to buy
1	ice-cream

Executor-2: F2

ID	Item to buy
2	chips

...

Executor-1000: F-1000

ID	Item to buy
1000	grapes

How do we write parallel programs?

- **Task parallelism**

- Partition various tasks carried out solving the problem among the cores.

- **Data parallelism**

- Partition the data used in solving the problem among the cores.
- Each core carries out similar operations on it's part of the data.

Data parallelism

- Let your data has 200,000,000,000 data points
- Partition your data into 100,000 chunks:
 - Number of partitions: **100,000**
 - Number of records per partition: **2000,000**
 - **100,000 x 2000,000 = 200,000,000,000**
- Assume you want to execute **map(function)** on (a transformation function) each record and create a new record
- With these partitioning in place, The fastest way to execute **map(function)** will be to have **100,000** mappers (mappers are transformations, which execute in parallel)
- **What happens if we have only 1000 mappers?**

Data parallelism

- What happens if we have only 1000 mappers?
- First, we assign 1000 of these partitions to 1000 mappers (**each mapper gets a single partition, which has 2,000,000 records**)
- Once a mapper completes its task, we assign another partition to that mapper
- This iteration continues until we exhaust all 100,000 partitions
- **At most 1000 mappers are executing at a single point of time.**
- **The more mappers we have:**
The more we execute faster

Parallelism requires Coordination

- Executors usually need to coordinate their work.
- **Communication** – one or more Executors send their current partial results to another core.
- **Load balancing** – share the work evenly among the Executors so that one is not heavily loaded.
- **Synchronization** – because each Executors works at its own pace, make sure Executors do not get too far ahead of the rest.

NOTE: If you use Spark or MapReduce, then all of these are done automatically for you!!!



Partitioner: partition data into chunks

- A partitioner partitions the input into chunks.
- If input has 200,000,000,000 records and we partition this input into 200,000 chunks, then we have:
 - Number of partitions: 200,000
 - Size of each partition: 1000,000 records
 - $200,000,000,000 = 200,000 \times 1000,000$
- Typically, a chunk (1000,000 records) becomes a unit of parallelism.

Chunk = Partition



Benefits of Parallel Computing

- **Parallel computing models the real world.** The world around us isn't serial and sequential: many things happen at the same time
- **Saves time.** Serial/sequential computing forces fast processors to do things inefficiently.
- **Saves money.** By saving time, parallel computing makes things cheaper and faster
- **Solve more complex or larger problems** by partitioning them into smaller problems
- **Solve Larger Problems in a short point of time.**



References

- **Introduction to Parallel Computing**
 - <https://www.geeksforgeeks.org/introduction-to-parallel-computing/>
- **Introduction: Parallelism = Opportunities + Challenges**
 - <https://courses.cs.washington.edu/courses/csep524/07sp/poppChaper1.pdf>

